

GPGPU Performance Modeling using Microbenchmarks



Hyounghoo Kim (kimhj1473@snu.ac.kr), Professor: Jangwoo Kim, TA: Joonsung Kim

Introduction

- ✓ GPU **internal structure is important**
 - Architecture research
 - Application performance
 - Simulator accuracy
- ✓ Understanding such **detail is difficult**
 - Manufacturer secret
 - Complexity
 - Rapidly changing
- **Systematic, Generic GPU Reverse-Engineering Tool?**

Background

GPU Programming Model


- Kernel: GPU code
 - B threads in a block
 - G blocks in a grid
- Warp: SIMT unit
 - w threads in a warp
- MultiProcessor (Indep.)

Here, we define
 $b = \lceil B/w \rceil$, $g = \lceil G/n_{mp} \rceil$

Related Works

- vs. Microbenchmark works
 - Also focused on HW resources
 - Portable among Nvidia, AMD
 - Warp scheduling policy
- vs. Modeling works
 - Parametrized
 - Models HW, not workload
 - Functional unit formula

GPGPU Modeling



- **Limit Checker:** Checks kernel size
 $B \leq L_{TB}, BG \leq L_{TG}, S \leq L_{SB}, r_i \leq L_{R_iT}, br_i \leq L_{R_iB}$
- **Block Scheduler:** Allocate blocks
 $T(g, b) = \lfloor \frac{g}{N_{slot}} \rfloor fu(bN_{slot}) + fu(b(g\%N_{slot}))$
- **State Buffers:** Warp slots, RF, sharedM
 $N_{slot} = \min(N_1, \lfloor \frac{N_2}{b} \rfloor, \lfloor \frac{S}{s_{min}} \rfloor, \lfloor \frac{R_i}{r_{i,min}b} \rfloor)$
- **Functional Units:** Alloc & exec warps
 $fu(c) = \max(1, \lfloor \frac{\sum u}{m_u p} \rfloor \lfloor \frac{c}{s_u} \rfloor | u \in FU)$
- **Cache Hierarchy**

Microbenchmarks

Limit Checker

Try & Find Max

Functional Units

```
measure_fu(G, B):
    thread_barrier()
    start = clock()
    #unroll rep(N):
        asm(instruction)
    thread_barrier()
    return clock() - start
```

(G,B) = (1,1): Latency
 (1,w..bw): Hardware Width
 (1,n), inst=clock(): Sched. Policy

State Buffers

```
measure_sb(G, B, *sync):
    if(tid=0) atomicAdd(sync, 1)
    while(*sync < G)
        check_if_timeout()
    return
    use_resources(s, r_i)
```

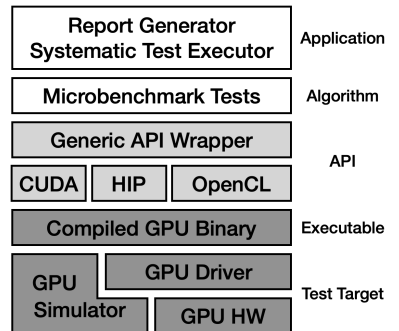
This **deadlocks** if $\frac{G}{n_{mp}} > N_{slot}$

(G,B) = (1..,c): Number of MP
 (1..,1..): Warp state buffer size
 (1..,1..), s=1..: SharedM size
 (1..,1..), r=1..: RF size

Cache Hierarchy

Fine-grained p-chase, portable

Implementation




Verification Method


- **Tested environments**
 - (A) gpgpusim (Turing, PTX, CUDA)
 - (B) gem5-apu (gfx803, GCN3, HIP)
 - (C) AMD V520 (gfx1011, RDNA1, HIP)
- **Parameter value comparison**
 Simulators: Config values and source code
 Hardware: Official spec, if available

Verification Result

- **Simple discrete parameters**
 Limits, number of MPs, cache linesize, ...
Measured accurately
- **Functional unit width**
Match with the model, little error




- **State buffer size**
Match with the model, some differs in (C)



Case Studies


Warp scheduling policy



(gem5 Oldest First) (AMD V520)


- gem5 **can't simulate** the sched policy of **real AMD hardware**, which is not RR
- Right graph's shape is similar to gpgpusim Greedy Than Oldest policy

Block allocation inefficiency



- For simple reg alloc policy in gem5, $b \geq 2$ kernel execution time at $g=4n+3, 4$ is **larger than expected**
- Due to the gem5 code nature
- Should use dynamic policy for accurate simulation

RDNA warp state buffer



- Differ slightly from model
- Total warp slot number is **not multiple of nmp**
- may be asymmetric?
- gem5 **needs other thread slot model** to support RDNA

Conclusion

- ✓ GPU Model
 - ✓ Generic
 - ✓ Parametrized
- ✓ Microbenchmarks
- ✓ Implemented
 - ✓ One-Click Runnable
- ✓ Verified
 - ✓ Simulator & HW
- ✓ Found gem5 faults
 - Anomalies in FU data
 - + Can go deeper
 - + RDNA support
 - + Intel, ARM GPUs?